

# Simple Convolutional Neural Network From Scratch

Chris Yogodzinski

## Notebook Overview

This notebook is a follow up to `src/python/scratch_two_layer.ipynb`.

The previous notebook I made a simple neural network from scratch to predict handwritten digits in the mnist dataset. The neural network consisted of one fully connected hidden layer and an output layer.

In this notebook I will be adding a convolutional layer and pooling layer to this simple neural network. The convolutions in this neural network function to key features from the dataset prior to the predictive fully connected layer.

## Variables

$I$  matrix: a mnist image of dim = ( $m$  by  $n$ )

$K$  matrix: kernel of dim = ( $m_k$  by  $n_k$ )

$s$  integer: stride, a constant set by user

$\phi_a$  function: activation function of convolution

$F$  matrix: a feature map of dim = ( $m - m_k + 1$  by  $n - n_k + 1$ ) or ( $i$  by  $j$ )

$\phi_p$  function: pooling function applied to feature map

$P$  matrix: pooled feature map, condensed image of dim = ( $\frac{m-m_k}{s}$  by  $\frac{n-n_k}{s}$ )

## Convolutional Layer

The convolutional layer in a cnn is named after a method known as kernel convolution.

Kernel convolution works by passing a filter or kernel iteratively over the input matrix and transform the values of the input based on the values of the kernel.

In this notebook our kernel transformation will be  $\sum_i \sum_j I[m-i, n-j] \times K$ .

In words this is element-wise multiplication of the kernel-sized subset of  $I$  and  $K$  followed by the sum of the product.

```
setwd("~/Projects/nn_playground")

read_mnist <- function(label_path, im_path) {
  # read in labels
  f <- file(label_path, "rb")
  meta <- readBin(f, n = 2, "integer", endian = "big")
  labels <- as.integer(readBin(f, n = meta[2], "raw", endian = "big"))
  close(f)

  # read in imgs
  f <- file(im_path, "rb")
  meta <- readBin(f, n = 4, "integer", endian = "big")
  byte_count <- meta[2] * meta[3] * meta[4]
```

```

    imgs <- readBin(f, n = byte_count, "raw", endian = "big")
    close(f)
    imgs <- array(as.integer(imgs), dim = c(meta[3], meta[4], meta[2]))
    dat <- list(labels = labels, imgs = imgs)
    return(dat)
}

dat_path = "data/mnist"
train_dat <- read_mnist(
  paste(dat_path, "train-labels.idx1-ubyte", sep = "/"),
  paste(dat_path, "train-images.idx3-ubyte", sep = "/")
)

test_dat <- read_mnist(
  paste(dat_path, "t10k-labels.idx1-ubyte", sep = "/"),
  paste(dat_path, "t10k-images.idx3-ubyte", sep = "/")
)

```

## Flipping The Images

The function above reads the images in upside down for some reason. In the code chunk below I flip the images right side up and visualize them.

```

flip_all_imgs <- function(array) {
  flip_image <- function(img) {
    flipped <- img[, ncol(img):1]
    return(flipped)
  }
  dims <- dim(array)
  flipped <- array(apply(array, 3, flip_image), dim = dims)
  return(flipped)
}

train_dat[["imgs"]] <- flip_all_imgs(train_dat[["imgs"]])
test_dat[["imgs"]] <- flip_all_imgs(test_dat[["imgs"]])


im_num <- 5

# test image
img <- test_dat[["imgs"]][ , , im_num]
cat(test_dat[["labels"]][im_num])

```

```
## 4
```

```
image(img, col = gray((0:255) / 255))
```



```
# train image
img <- train_dat[["imgs"]][ , , im_num]
cat(train_dat[["labels"]][im_num])
```

```
## 9
```

```
image(img, col = gray((0:255) / 255))
```

